

Adversarial Search

Games vs. search problems

- "Unpredictable" opponent → specifying a move for every possible opponent reply
-
- Time limits → unlikely to find goal, must approximate
-


Game search

- Game-playing programs developed by AI researchers since the beginning of the modern AI era
 - Programs playing chess, checkers, etc (1950s)
- Specifics:
 - Sequences of player's decisions we control
 - Decisions of other player(s) we do not control
- Contingency problem: many possible opponent's moves must be “covered” by the solution
- Opponent's behavior introduces uncertainty
- Rational opponent – maximizes its own utility (payoff) function

Game Search Problem

- Problem formulation
 - Initial state: initial board position + whose move it is
 - Operators: legal moves a player can make
 - Goal (terminal test): game over?
 - Utility (payoff) function: measures the outcome of the game and its desirability
- Search objective:
 - Find the sequence of player's decisions (moves) maximizing its utility (payoff)
 - Consider the opponent's moves and their utility

Game tree (2-player, deterministic, turns)



Minimax Algorithm


- How to deal with the contingency problem?
 - Assuming the opponent is always rational and always optimizes its behavior (opposite to us), we consider the best opponent's response
 - Then the minimax algorithm determines the best move

Minimax


- Perfect play for deterministic games
-
- Idea: choose move to position with highest **minimax value**
= best achievable payoff against best play
-
- E.g., 2-ply game: **[will go through another eg in lecture]**

MAX

MIN



Minimax. Example



function MINIMAX-DECISION(*game*) **returns** *an operator*

for each *op* **in** OPERATORS[*game*] **do**

 VALUE[*op*] \leftarrow MINIMAX-VALUE(APPLY(*op*, *game*), *game*)

end

return the *op* with the highest VALUE[*op*]

function MINIMAX-VALUE(*state*, *game*) **returns** *a utility value*

if TERMINAL-TEST[*game*](*state*) **then**

return UTILITY[*game*](*state*)

else if MAX is to move in *state* **then**


return the highest MINIMAX-VALUE of SUCCESSORS(*state*)

else

return the lowest MINIMAX-VALUE of SUCCESSORS(*state*)

Complexity of the minimax algorithm

- We need to explore the complete game tree before making the decision



- Impossible for large games
 - Chess: 35 operators, game can have 50 or more moves

Solution to the complexity problem

Two solutions:

1. **Dynamic pruning of redundant branches** of the search tree

- identify a provably suboptimal branch of the search tree before it is fully explored
- Eliminate the suboptimal branch

Procedure: Alpha-Beta pruning


2. **Early cutoff of the search tree**

- uses imperfect minimax value estimate of non-terminal states (positions)

Alpha Beta Pruning

- Some branches will never be played by rational players since they include sub-optimal decisions for either player
- First, we will see the idea of Alpha Beta Pruning
- Then, we'll introduce the algorithm for minimax with alpha beta pruning, and go through the example again, showing the book-keeping it does as it goes along


Alpha beta pruning. Example



Minimax with alpha-beta pruning: The algorithm

- Maxv: function called for max nodes
 - Might update alpha, the best max can do far
- Minv: function called for min nodes
 - Might update beta, the best min can do so far
- Each tests for the appropriate pruning case
- We'll go through the algorithm on the course website

Algorithm example: alphas/betas shown




Properties of α - β

- Pruning **does not** affect final result
-
- Good move ordering improves effectiveness of pruning
-
- With "perfect ordering," time complexity = $O(b^{m/2})$
 - **doubles** depth of search
- A simple example of the value of reasoning about which computations are relevant (a form of **metareasoning**)
-

Using minimax value estimates

- **Idea:**
 - Cutoff the search tree before the terminal state is reached
 - Use imperfect estimate of the minimax value at the leaves
 - Evaluation function



Design of evaluation functions

- **Heuristic estimate** of the value for a sub-tree
- **Examples of a heuristic functions:**
 - **Material advantage in chess, checkers**
 - Gives a value to every piece on the board, its position and combines them
 - More general **feature-based evaluation function**
 - Typically a linear evaluation function:

$$f(s) = f_1(s)w_1 + f_2(s)w_2 + \dots + f_k(s)w_k$$

$f_i(s)$ - a feature of a state s

w_i - feature weight

Further extensions to real games

- Restricted set of moves to be considered under **the cutoff level** to reduce branching and improve the evaluation function
 - E.g., consider only the capture moves in chess

